


# ECE 437 / CS 481 Operating Systems

---

By  
Dr. Edward Nava

# Course objectives

---



- To learn “what an operating system is and how it works”
- To understand OS design principles
- To understand “how OS concepts are implemented”
- To practice systems programming skills
- To learn about OS performance and its evaluation

***A. Silberschatz, Peter Baer Galvin and Greg Gagne  
" Operating System Concepts", 9<sup>th</sup> edition, ISBN  
978-1-118-06333-0, Wiley, Inc, 2013.***

# Course Details

---

- Lecture
  - DSH 120
  - Mon/Wed 11:00-12:15pm
- Prerequisites
  - ECE331 (Data structures & C Programming)
  - ECE344L or CS341 (Computer Organization background)
- Office Hours and Location
  - ECE 225C
  - MW 1:30-2:30pm or by emailing for appointment
  - ejnava@unm.edu

# Course Outline

---

- OS overview, Linux & Windows
- Processes & Threads
- CPU scheduling
- Process synchronization
- Deadlocks
- Files and file systems
- I/O devices & management
- Memory management
- Protection & security



# Course Grading

---

- Homework Assignment and Programming Projects: 40%
  - Homework and projects are intended to be individual assignments unless team work is specifically permitted. Evidence of copying will be penalized; all guilty parties will be given a zero.
  - Late programming assignments will be accepted within 3 calendar days but with a 50% PENALTY.
  - No late homework will be accepted; No makeup quizzes;
- Exams 60%
  - Exams will cover material presented in class, programming projects, and homework assignments.
  - Final Exam (optional, comprehensive): Replace min(Exam I, II, III)

# Course logistics

---

- Course GA or PA
  - Mithun Mohan, mithunmohan@unm.edu
- Check course homepage for assignment/grades
  - <http://learn.unm.edu>
- Tentative schedule is posted at learn.unm.edu and will be updated through the semester
- Tentative exam dates
  - Exam I, Tuesday, Aug 26
  - Exam II, Tuesday, Oct 31
  - Exam III, Thursday, Nov 30
  - Final Exam week of Dec 5

# Lecture 1: Introduction



# Lecture 1: Introduction

---

- What Operating Systems Do
- Computer-System Organization





# What is an Operating System?

---


- A program that acts as an intermediary between a user of a computer and the computer hardware
  - Operating system goals:
    - Execute user programs and make solving user problems easier
    - Make the computer system convenient to use
    - Use the computer hardware in an efficient manner
-



# Computer System Organization

---

- Computer-system operation
  - One or more CPUs, device controllers connect through common bus providing access to shared memory
  - Concurrent execution of CPUs and devices competing for memory cycles



# Computer Startup

---

- **bootstrap program** is loaded at power-up or reboot
  - Typically stored in ROM or EPROM, generally known as **firmware**
  - Initializes all aspects of system
  - Loads operating system kernel and starts execution



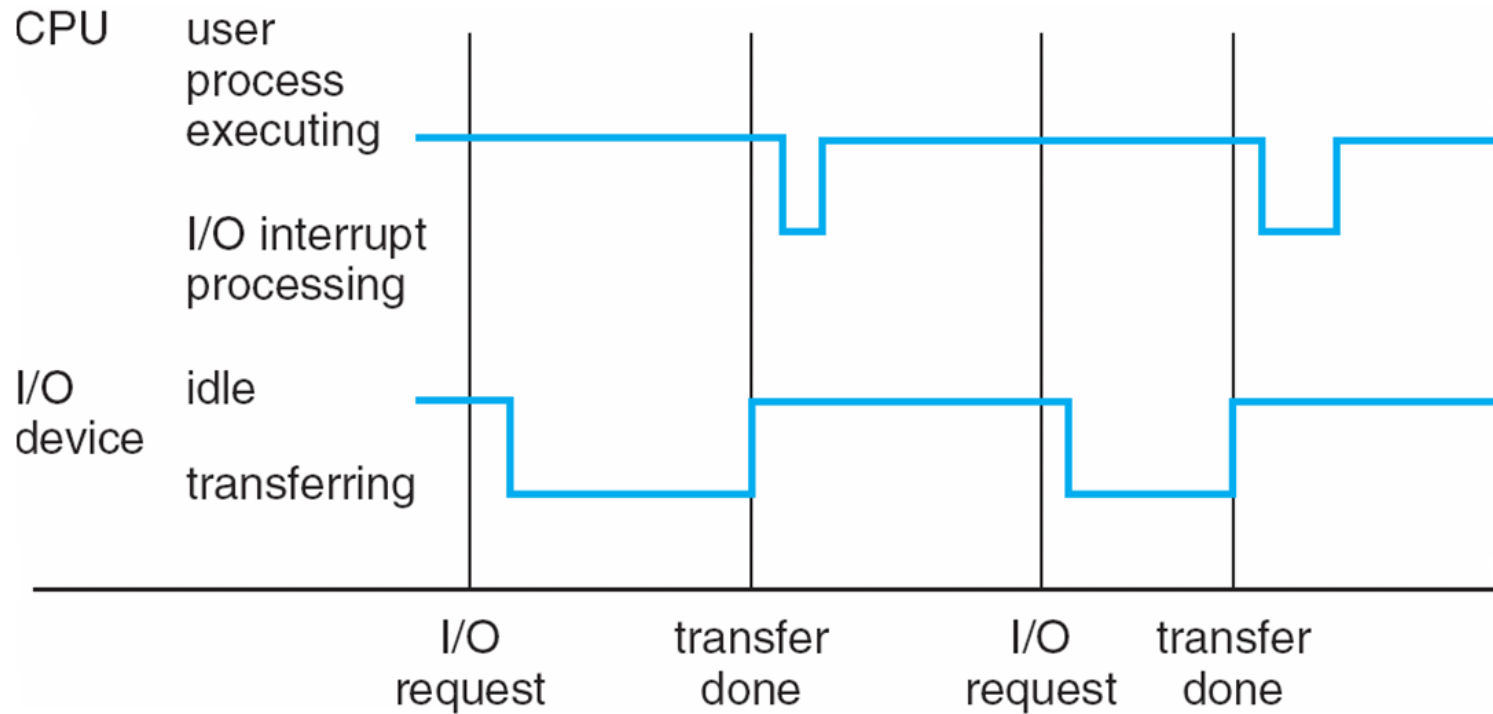
# Common Functions of Interrupts

---

- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
  - Interrupt architecture must save the address of the interrupted instruction
  - Incoming interrupts are *disabled* while another interrupt is being processed to prevent a *lost interrupt*
  - A *trap* is a software-generated interrupt caused either by an error or a user request
  - An operating system is **interrupt driven**
-

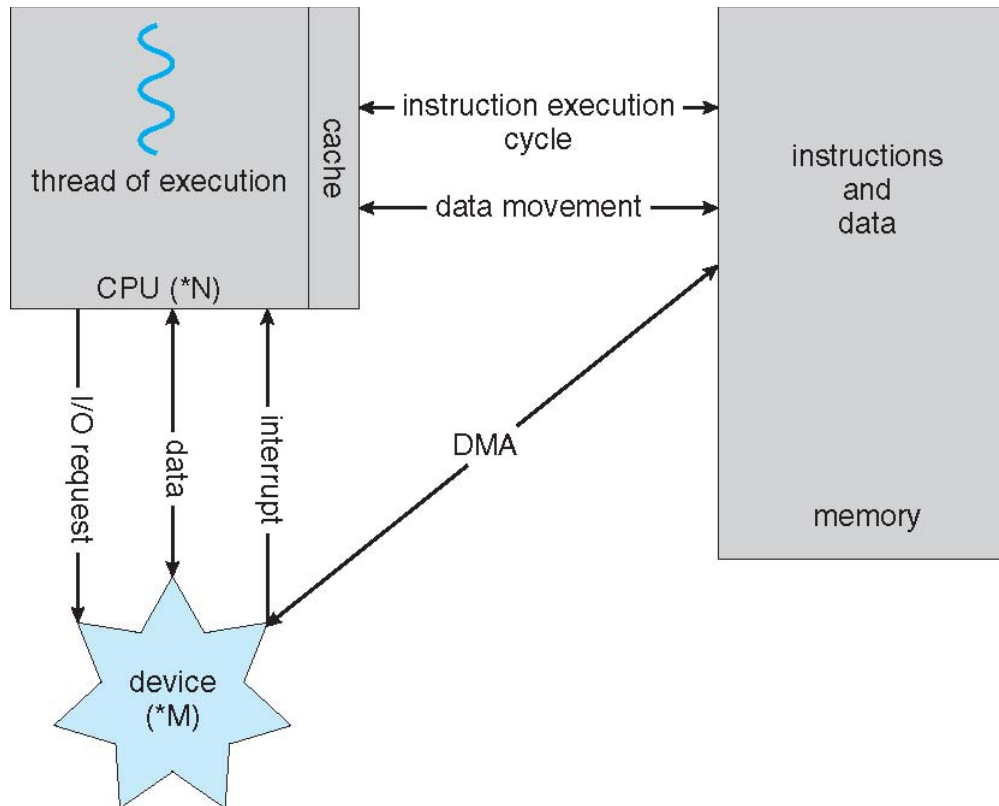


# Interrupt Timeline



# How a Modern Computer Works

---





# Computer-System Operation


---

- I/O devices and the CPU can execute concurrently
  - Each device controller is in charge of a particular device type
  - Each device controller has a local buffer
  - CPU moves data from/to main memory to/from local buffers
  - I/O is from the device to local buffer of controller
  - Device controller informs CPU that it has finished its operation by causing an *interrupt*
-



# Computer System Organization

- Computer-system operation
  - One or more CPUs, device controllers connect through common bus providing access to shared memory
  - Concurrent execution of CPUs and devices competing for memory cycles





# Storage Definitions and Notation Review

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.

A **kilobyte**, or **KB**, is 1,024 bytes

a **megabyte**, or **MB**, is  $1,024^2$  bytes

a **gigabyte**, or **GB**, is  $1,024^3$  bytes

a **terabyte**, or **TB**, is  $1,024^4$  bytes

a **petabyte**, or **PB**, is  $1,024^5$  bytes

Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).



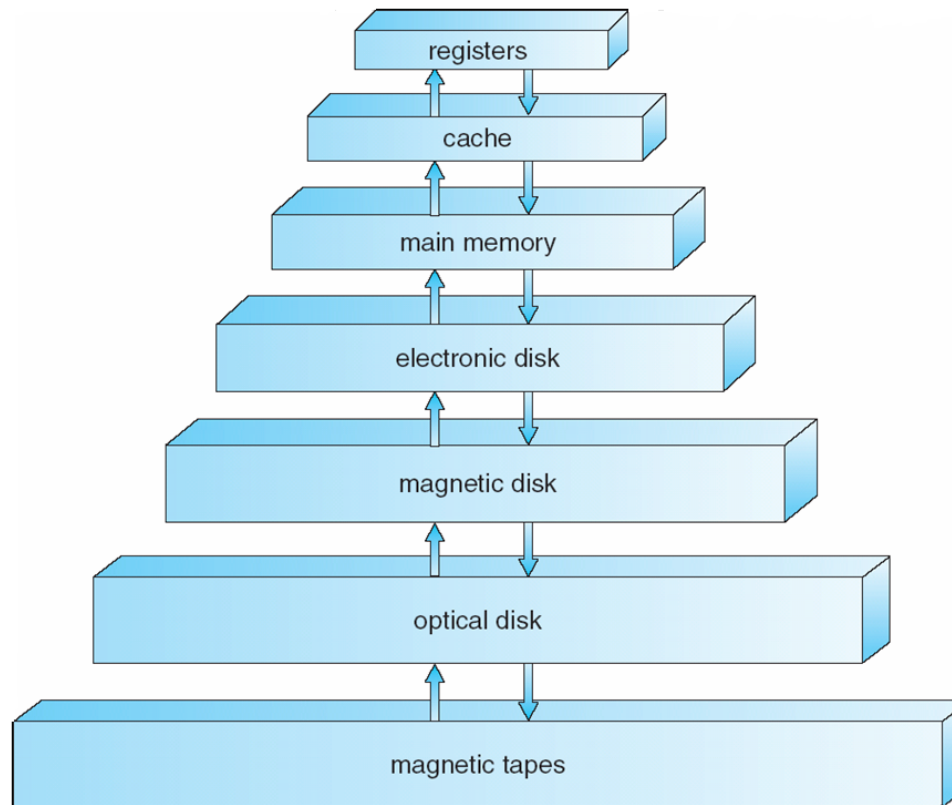
# Storage Structure

---

- Main memory – only large storage media that the CPU can access directly
    - **Random access**
    - Typically **volatile**
  - Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
  - Hard disks – rigid metal or glass platters covered with magnetic recording material
    - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
    - The **disk controller** determines the logical interaction between the device and the computer
  - **Solid-state disks** – faster than hard disks, nonvolatile
    - Various technologies
    - Becoming more popular
-

# Storage-Device Hierarchy

---





# Caching

---

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
  - Information in use copied from slower to faster storage temporarily
  - Faster storage (cache) checked first to determine if information is there
    - If it is, information used directly from the cache (fast)
    - If not, data copied to cache and used there
  - Cache smaller than storage being cached
    - Cache management important design problem
    - Cache size and replacement policy
-



# Direct Memory Access Structure

---

- Used for high-speed I/O devices able to transmit information at close to memory speeds
  - Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
  - Only one interrupt is generated per block, rather than the one interrupt per byte
-